

6.854
ADVANCED ALGORITHMS
PROBLEM SET 1

JOHN WANG

Collaborators: Jason Hoch, Ryan Liu, Varun Ganesan

1. PROBLEM 1

Problem: Unlike regular heaps, Fibonacci heaps do not achieve their good performance by keeping the depth of the heap small. Demonstrate this by exhibiting a sequence of Fibonacci heap operations on n items that produce a heap-ordered tree of depth $\Omega(n)$.

Solution: Consider the following recursive series of operations. For the i th step of the recursion, we will assume there is a tree t_1 of depth i , composed of exactly i nodes. There is also a tree t_2 which is a single node such that $root(t_1) < root(t_2)$. We shall insert two nodes a and b such that $b < a < root(t_1) < root(t_2)$. Perform a delete-min operation on the set of trees. This operation will remove b since it is the minimum of the entire structure. This leaves us with 3 roots, namely $a, root(t_1)$, and $root(t_2)$. The delete-min operation will also perform a consolidation, so that a is merged with t_1 , then the resulting tree is merged with t_2 .

The resulting tree has a root of a , a left child of $root(t_1)$ and a right child of $root(t_2)$. Now, we perform a decrease key on $root(t_2)$ to a value lower than a . This will cut it off from the tree, and we will be left with a tree t'_1 rooted at a and t'_2 . We see that t'_1 will have a depth of $i + 1$ and t'_2 will be a single node. Thus, we are back to our original datastructure with step $i + 1$ and can recurse.

Note that we performed four operations: insert a , insert b , delete-min, decrease-key. Thus, we see that after n of these operations, we will have a tree of length $n/4 = \Omega(n)$. \square

2. PROBLEM 2

Problem: Suppose that Fibonacci heaps were modified so that a node was cut only after losing k children. Show that this will improve the amortized cost of decrease key (to a better constant) at the cost of a worse cost for delete-min (by a constant factor).

Solution: First we will define our potential function as $\Phi = R + 2M/(k - 1)$ where R is the number of roots and M is the number of mark bits. Examining the amortized cost of insert a_i is given by:

$$(1) \quad a_i = c + 1 + \Delta\Phi$$

Where c is the number of nodes cut on a given insert (due to cascading). The real cost is $c + 1$ because c nodes are cut during cascading, each requiring constant time, and $+1$ because the node must be inserted into the data structure as well. The change in potential is given by:

$$(2) \quad \Delta\Phi = c + \frac{2(1 - (k - 1)(c - 1))}{k - 1}$$

Because c nodes are cut during the cascading, an additional c roots are created which accounts for the first c term in $\Delta\Phi$. Moreover, the number of mark bits decreases by $(k - 1)(c - 1)$ since we cut away c nodes, which means that $c - 1$ of these nodes had $k - 1$ mark bits already stored which were cleared when everything was cascaded. However, we added 1 mark bit to the last node in the cascading chain, which is why we have a change of $1 - (k - 1)(c - 1)$ mark bits. Putting this into our expression, we obtain:

$$(3) \quad a_i = 1 + c + c + \frac{2(1 - (k - 1)(c - 1))}{k - 1}$$

$$(4) \quad = 1 + 2 + \frac{2}{k - 1}$$

$$(5) \quad = 3 + \frac{2}{k - 1}$$

Thus, when $k = 2$, the cost for insert is 5, whereas when $k > 2$, the cost for insert is less than 5. Thus, the change improves the amortized cost of decrease key to a better constant.

We are left to show that cutting nodes only after losing k children makes delete-min more expensive. We know that the real cost of delete min will be the number of roots r plus the maximum degree possible for the root. Since the potential function is only changed by the number of roots, we see that $\Delta\Phi = \max \text{degree} - r$. Therefore, we see that the amortized cost is simply $a_i = O(\# \text{ of roots})$.

Now to find the maximum degree, we need to consider node x with children y_0, y_1, \dots, y_i in the order they were added. We know y_i becomes a child of x during consolidation and the degree of x is at least $i - 1$ because y_0, \dots, y_{i-1} were already children of x . This means that y_i 's degree could only have decreased by $k - 1$ before it gets cut so that $\deg(y_i) \geq i - 1 - (k - 1) = i - k$. This means we can find a lower bound for the maximum degree, using S_m where S_m is the minimum number of nodes in a degree m subtree. This means:

$$(6) \quad S_m \geq \sum_{i=k}^m S_{i-k}$$

$$(7) \quad S_m - S_{m-1} \geq \sum_{i=k}^m S_{i-k} - \sum_{i=k}^{m-1} S_{i-k} = S_{m-k}$$

Since $S_m - S_{m-1} = S_{m-k}$, we find that characteristic polynomial can be written as $x^k - x^{k-1} = x^0$ or $x^k - x^{k-1} - 1 = 0$. When $k = 2$, we have $x^2 - x - 1 = 0$ so that $x_1 = \phi$. However, when $k > 2$, we have $x_2 < \phi$ so that $\log_{x_1}(n) < \log_{x_2}(n)$. Since we know that the maximum degree when $k = 2$ is less than the maximum degree when $k > 2$, we see that delete-min has worse constants in the case when $k > 2$. \square

3. PROBLEM 3

On tradeoffs in the heap operations.

Problem: Let P be a priority queue that performs insert, delete-min, and merge in $O(\log n)$ time, and performs make-heap in $O(n)$ time where n is the size of the resulting priority queue. Show that P can be modified to perform insert in $O(1)$ amortized time, without affecting the cost of delete-min or merge (i.e. $O(\log n)$ amortized time). Assume that the priority queue does not support an efficient decrease-key operation.

Solution: We will store a linked list L of priority queues. The original priority queue P will be at the front of the linked list. When performing an insert, we will create a new priority queue and append it to L . When merging a new priority queue P_{new} , we first consolidate all of the priority queues which are not P , hence all the priority queues which are single nodes and have been created by an insert, and create an auxiliary priority queue P_a . Next, we merge P_a with P so that the entire data structure is now composed of a single heap P . Then, we merge the new priority queue P_{new} with the existing queue P . To perform a delete-min, we first consolidate all of the roots into a single priority queue. This can be done with the same routine used in merge, i.e. creating a priority queue of the singular roots and merging that auxiliary priority queue with P . Then, delete min can be performed on P in the usual manner.

To analyze this data structure, we will introduce a potential function $\Phi = \#$ of roots. The real cost of an insert is just a constant $O(1)$, and the change in potential is 1, so $a_i = O(1)$ for insert. To examine merge and delete-min, we will first examine the consolidation subroutine. In the consolidation subroutine, a make-heap operation is performed on all single-node roots. The real cost of the make-heap is $O(c)$, where c is the number of single-node roots. The merge part of the consolidation requires $O(\log n)$ real cost. The change in potential is given by $-c$, since there are now c fewer roots. Thus, the amortized cost of the consolidation subroutine is $O(\log n) + c - c = O(\log n)$. This means that the merge operation, which will merge another tree in $O(\log n)$ (and won't affect the potential), will cost $O(\log n) + O(\log n) = O(\log n)$ amortized time. The delete-min operation will perform a merge, then a regular delete-min from the priority queue P , which will take $O(\log n) + O(\log n) = O(\log n)$ amortized time.

Thus, we have created a data structure which allows insert in $O(1)$ but retains the amortized cost of delete-min and merge. \square

Problem: Using the above technique, show that even binary heaps can be modified to support insert in $O(1)$ amortized time while maintaining an $O(\log n)$ time bound for delete-min. Note that binary heaps do not support merge in $O(\log n)$ time.

Solution: We will start out with a linked list L of priority queues as before. When we insert a node, we will simply create a new root and append it to L . We will also keep a priority queue Q which holds the roots of all other priority queues, and at first it will only contain the root of P . To perform delete-min, we will first perform a make-heap on all of the single-node roots, and insert the new heap's root into Q . After

the new heap's root has been inserted into Q , we perform a delete-min on Q , then a delete-min on the tree which was the root of Q .

We will use a potential function $\Phi = \#$ of roots for the analysis. Insert requires amortized time $a_i = 1 + 1$, since it requires 1 in real cost to create a new node and the potential changes by 1. For delete-min, we first perform a make-heap on all of the single-node roots. Let's say that there exist c single-node roots, then the make heap requires $O(c)$ time. Inserting the heap into Q will require $O(\log n)$ time, since Q has a maximum size of n and therefore a maximum depth of $O(\log n)$. Delete-min on Q also requires $O(\log n)$. The next delete-min on the resulting heap requires $O(\log n)$ time it must also be of size $O(n)$, and hence height $O(\log n)$. The total amortized cost of delete-min is therefore $a_i = 3O(\log n) + c - c = O(\log n)$.

We have therefore achieved insert in $O(1)$ amortized time and delete-min in $O(\log n)$ time for a binary heap. \square

4. PROBLEM 5

Problem: Show how to use the techniques of persistent data structures to preprocess a tree in $O(n \log n)$ time so as to allow LCA queries to be answer in $O(\log n)$ time. Aim for a simple solution here, even if you solve part (b).

Solution: We will use a persistent version of the offline algorithm given in the problem. First, we associate with each node an extra field "name" and process the nodes of T in post order. However, we will make the tree into a persistent data structure by keeping a timestamped log in a fat node and using path copying for each union find we use.

Thus, as we perform union-find operations in the preprocessing stage, we will be using path compression and fat nodes to store previous changes in the "name" attribute. This means that find will take $O(\log t)$ to find the correct tree to use and an additional $O(\log n)$ to find the node in that particular tree. Thus, there is a total real cost of $O(\log t) + O(\log n)$ to perform a find operation. To construct the data structure, we will have a real cost of $O(\log n)$ for each find and union operation, since find takes $O(\log n)$ and dominates the union operation. This means that preprocessing costs $O(\log n)$ per node. Since there are a total of n find operations in the sequence, we will have $O(n \log n)$ cost for the preprocessing.

We will use the potential function $\Phi = \#$ live nodes. Thus, the amortized time of find will be $a_i = O(\log t) + O(\log n) = O(\log n)$ because $t = O(n)$ as only one operation is done for each node in the preprocessing phase. Thus, we see that find will require $O(\log n)$ time. Next, we know that the preprocessing time will require $O(n)$ times the cost of make set, union, and find on a node. Since each of these costs $O(\log n)$ time in a union-find data structure, we see that the preprocessing time is $O(n)$.

This means that we can perform LCA queries in $O(\log n)$ and preprocess the tree in $O(n \log n)$ time. \square